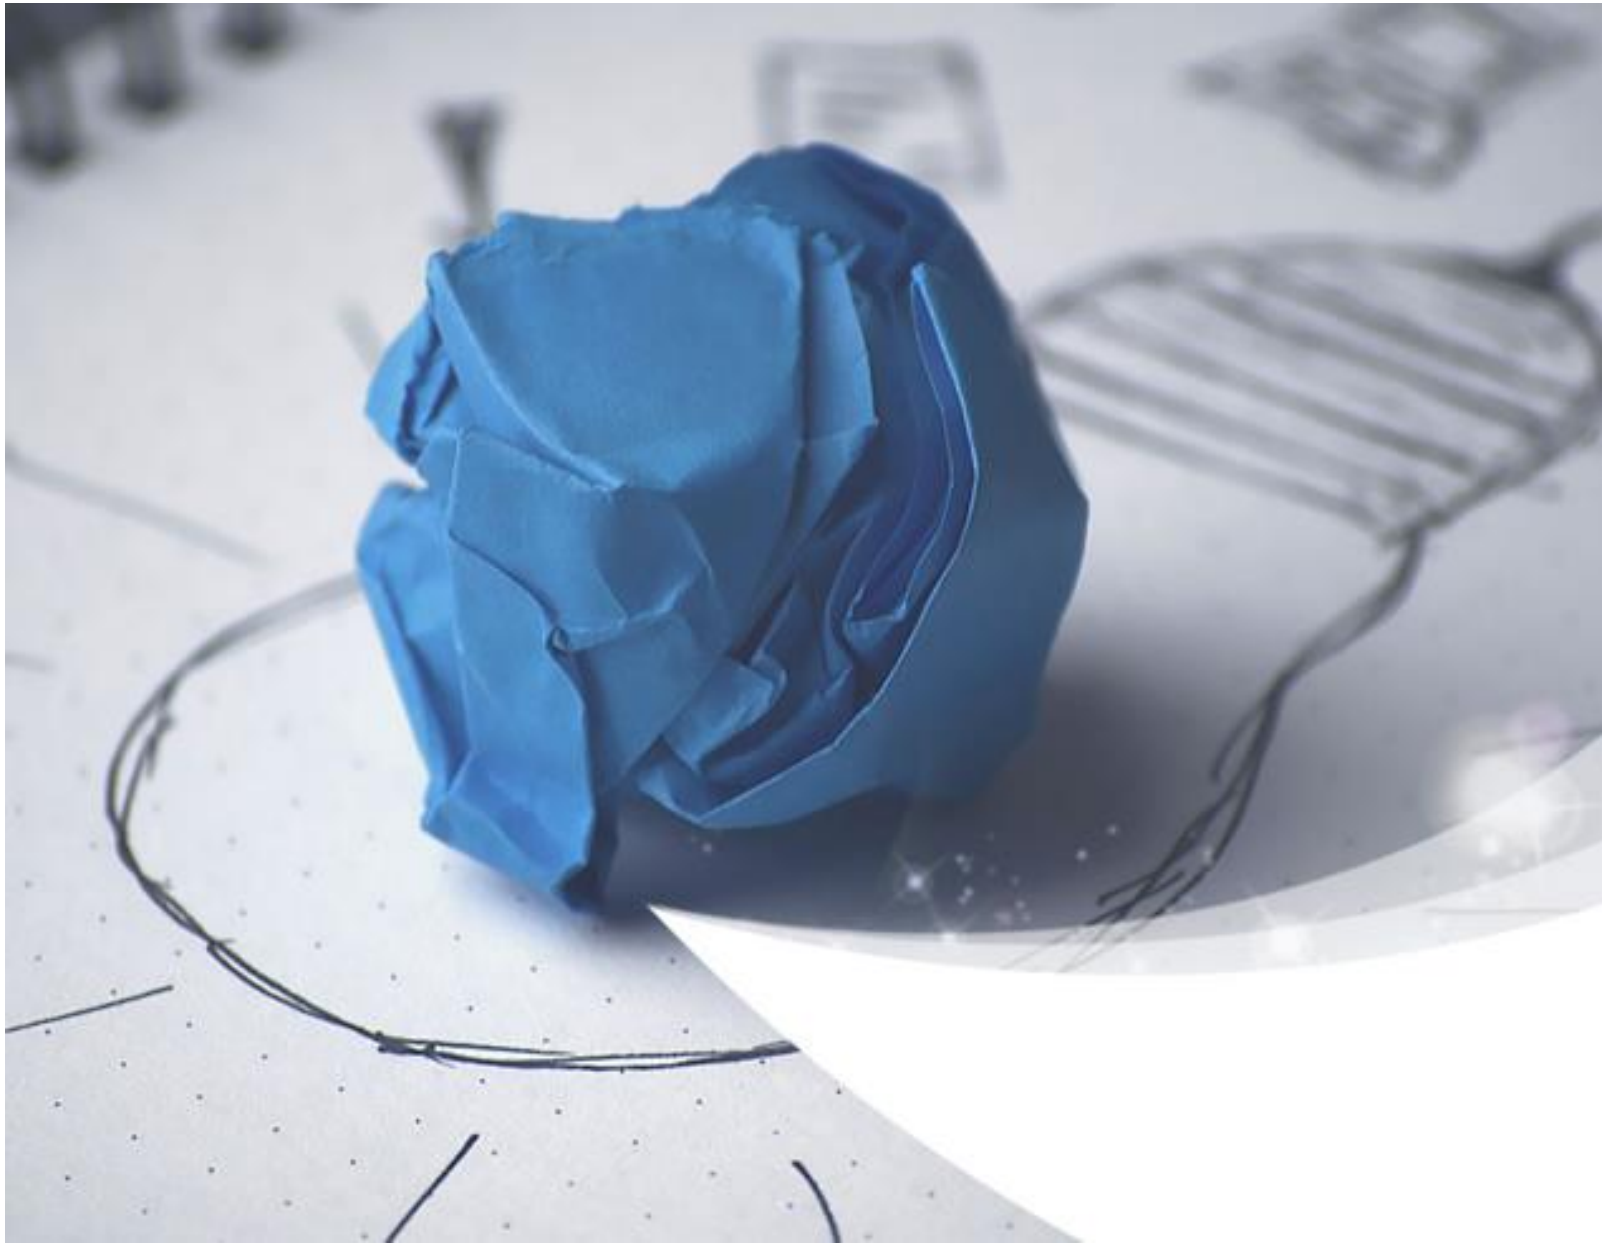

ADVANCED MACHINE LEARNING & TEXT MINING

Session 1

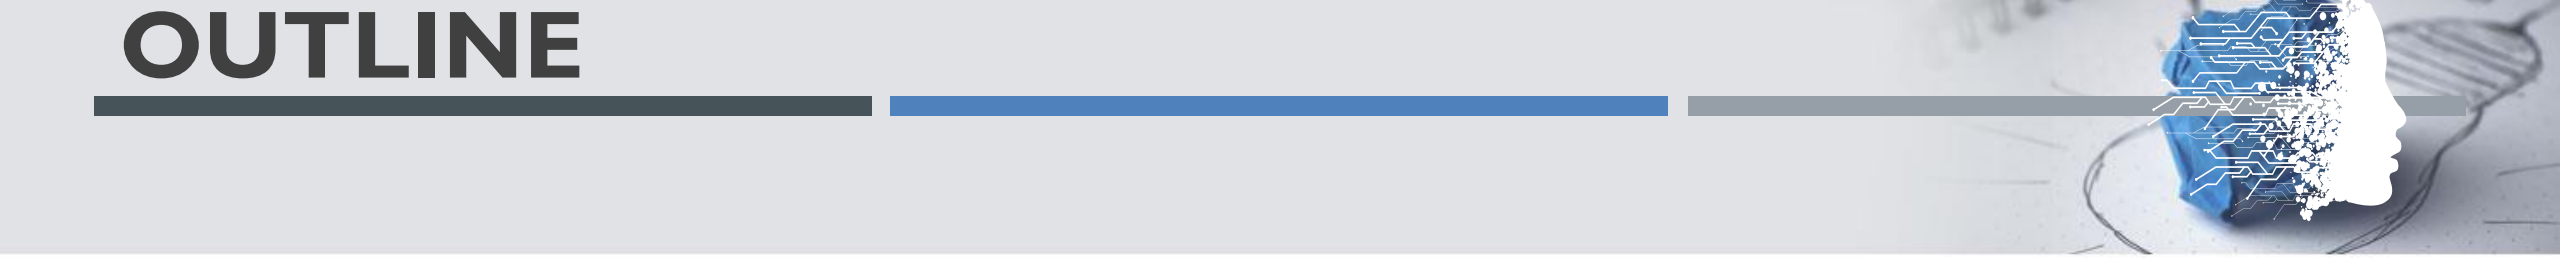

Data Pre-processing

TEXT MINING

- Definitions
  - Text mining, also referred to as text data mining, roughly equivalent to text analytics, refers to the process of deriving high-quality information from text." –Wikipedia
  - "Another way to view text data mining is as a process of exploratory data analysis that leads to heretofore unknown information, or to answers for questions for which the answer is not currently known." -Hearst, 1999
  - "Text mining applies techniques such as categorization, entity extraction, sentiment analysis and natural language processing to transform text into data that can be used for further analysis." – expert system

TEXT MINING

- Why Text Mining?
  - 1.8 Zettabytes unstructured data
  - Unlike search engines, which surface documents based on keywords, text mining tools analyze documents to identify entities and extract relationships between them, unlocking hidden information to help:
    - identify and develop new hypotheses attain knowledge
    - attain knowledge
    - improve understanding.

- Why Text Mining?
  - Text mining ensures the use of all of available information to make better informed decisions, automate information-intensive processes, gather business critical insights and mitigate operational risks.
  - Text mining can solve high-value knowledge discovery problems in many different areas of application.

- Types of text mining
  - **Search and information retrieval**
    - Storage and retrieval of text documents, including search engines, and keyword search
  - **Document Clustering**
    - Grouping and categorizing terms snippets, paragraph or documents using clustering algorithms
  - **Document classification**
    - Grouping and categorizing terms snippets, paragraph or documents using classification methods, based on models trained on labels example

- Types of text mining
  - **Web mining**
    - Data and text mining on the Internet with specific focus on scale and interconnectedness
  - **Information Extraction**
    - Identification and extraction relevant facts and relationships from unstructured text; the process of making structured data from unstructured and semi-structured data

- Types of text mining
  - **NLP**
    - Low level language processing and understanding tasks often used synonymously with computational linguistics
  - **Concept extraction**
    - Grouping of words and phrases into semantically similar groups

- Text mining around us

Sentiment Analysis

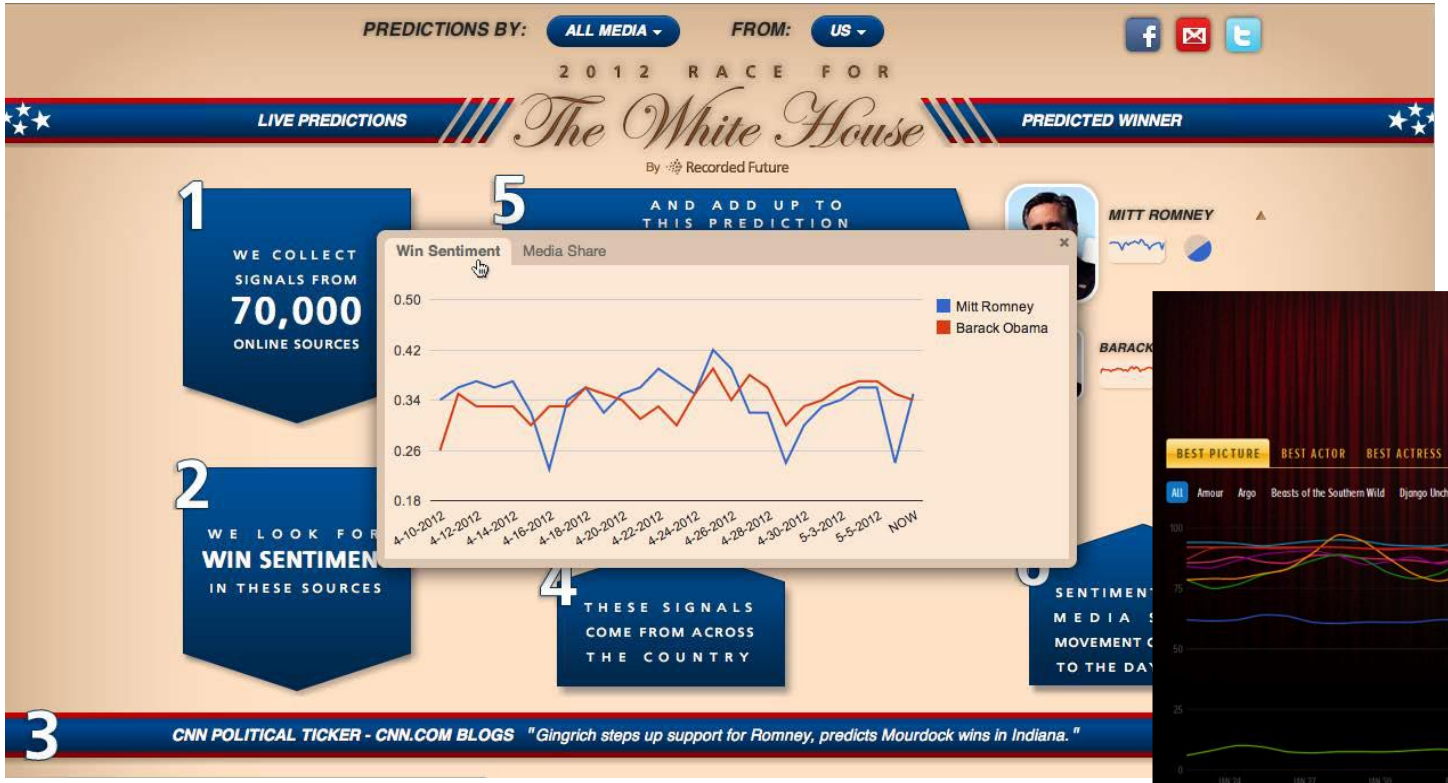

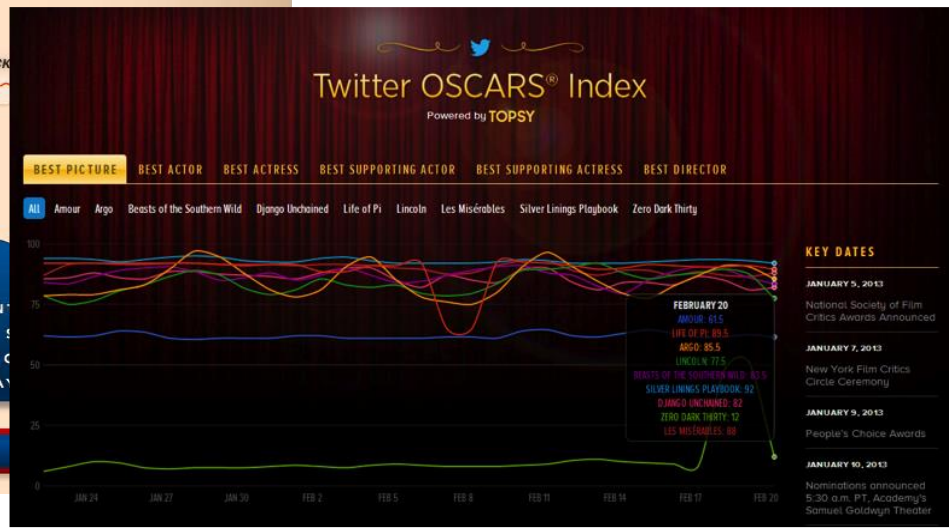

- Text mining around us

Document Summarization

- Text mining around us

Document Summarization

bing text mining

Web Images Videos Maps News More

19,200,000 RESULTS Any time

Text mining - Wikipedia, the free encyclopedia
en.wikipedia.org/wiki/Text_mining
Text mining, also referred to as text data mining, roughly equivalent to text analytics, refers to the process of deriving high-quality information from text. High ...
Text mining and text ... · History · Text analysis processes · Applications

Text Mining (Big Data, Unstructured Data)
www.statsoft.com/Textbook/Text-Mining
Text Mining Introductory Overview. The purpose of Text Mining is to process unstructured (textual) information, extract meaningful numeric indices from the text, ...

Text Mining
academic.research.microsoft.com/Keyword/41731/text-mining
Text mining is defined as knowledge discovery in large text collections. It detects interesting patterns such as clusters, associations, deviations, similarities, and ...

What is **text mining** (**text** analytics)? - Definition from ...
searchbusinessanalytics.techtarget.com/definition/text-mining
Text mining is the analysis of data contained in natural language text. The application of text mining techniques to solve business problems is called text analytics.

Text mining

Text mining, also referred to as text data mining, roughly equivalent to text analytics, refers to the process of deriving high-quality information from text. High-quality information is typically derived through the devising of patterns and trends through means such as statistical pattern learning. Text mining usually involves the process of struct... +

en.wikipedia.org

Related people: Junichi Tsujii · Alfonso Valencia · Tomoko Ohta · Carol Friedman · Michael Berry · Hsinchun Chen

People also search for: Sentiment analysis · Natural language processing · Web mining · Analytics · Cluster analysis +

Data from: Wikipedia · Freebase

Feedback

Related searches
Text Analysis Software
Text Analytics

- Text mining around us

Movie Recommendation

- Text mining around us

Text mining in healthcare

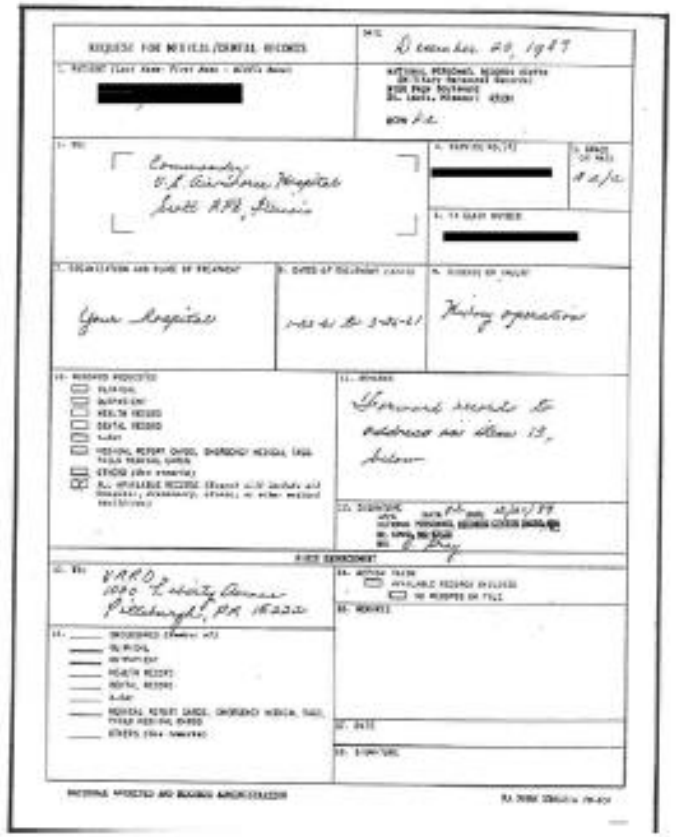

- Text mining around us

Text mining in Finance

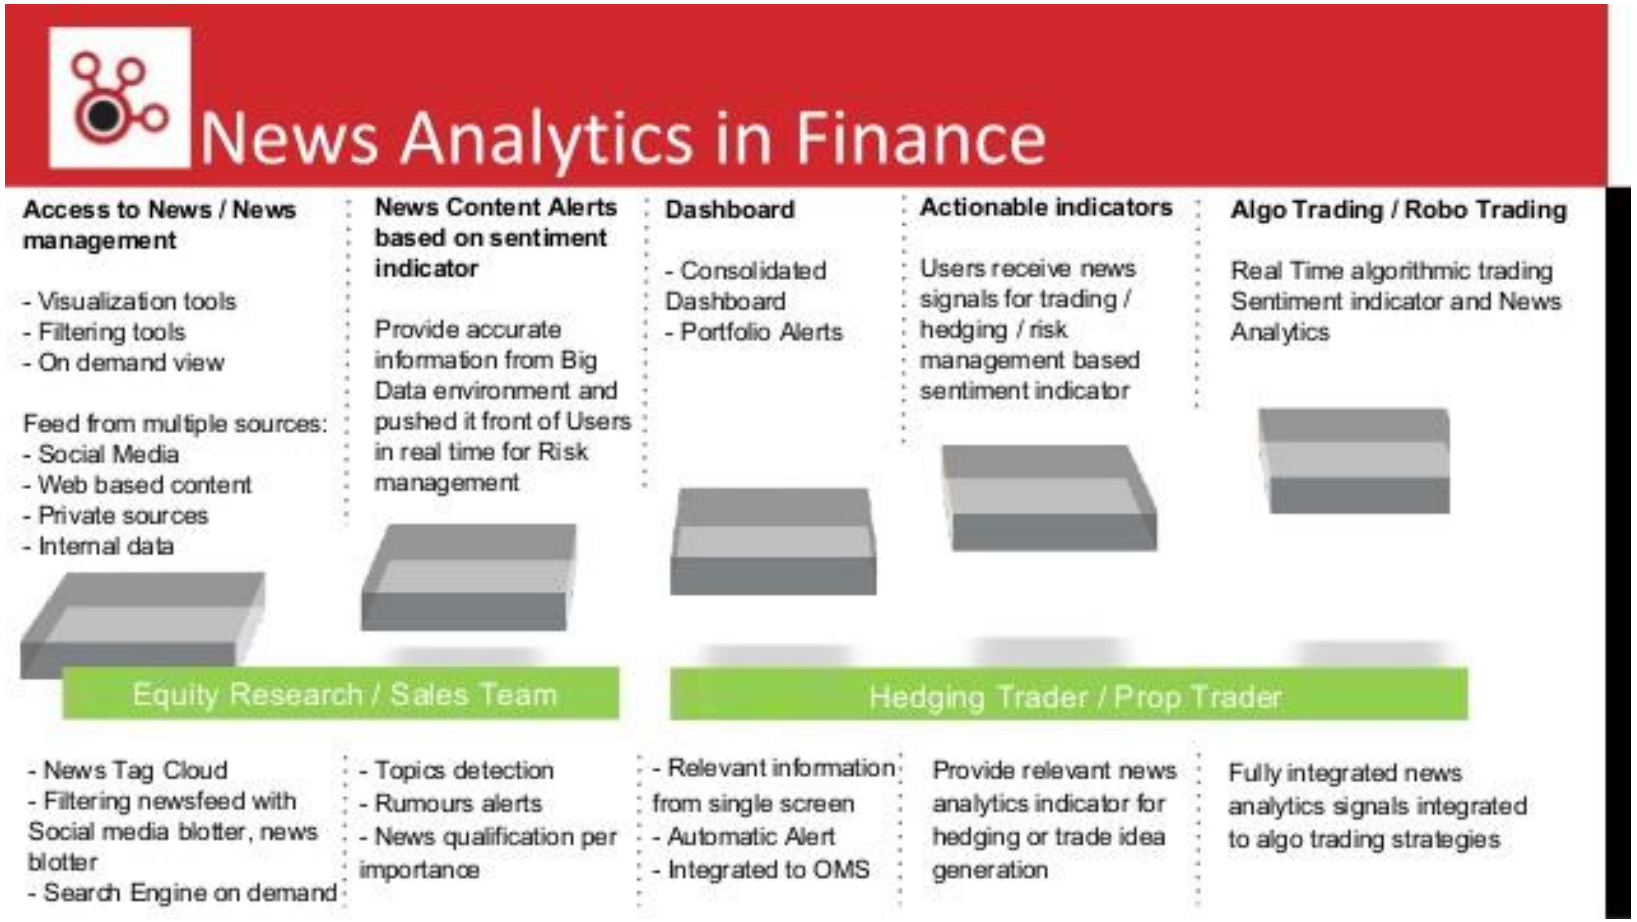

- Text Mining Process Flow

(Source: Claudia Peersman)

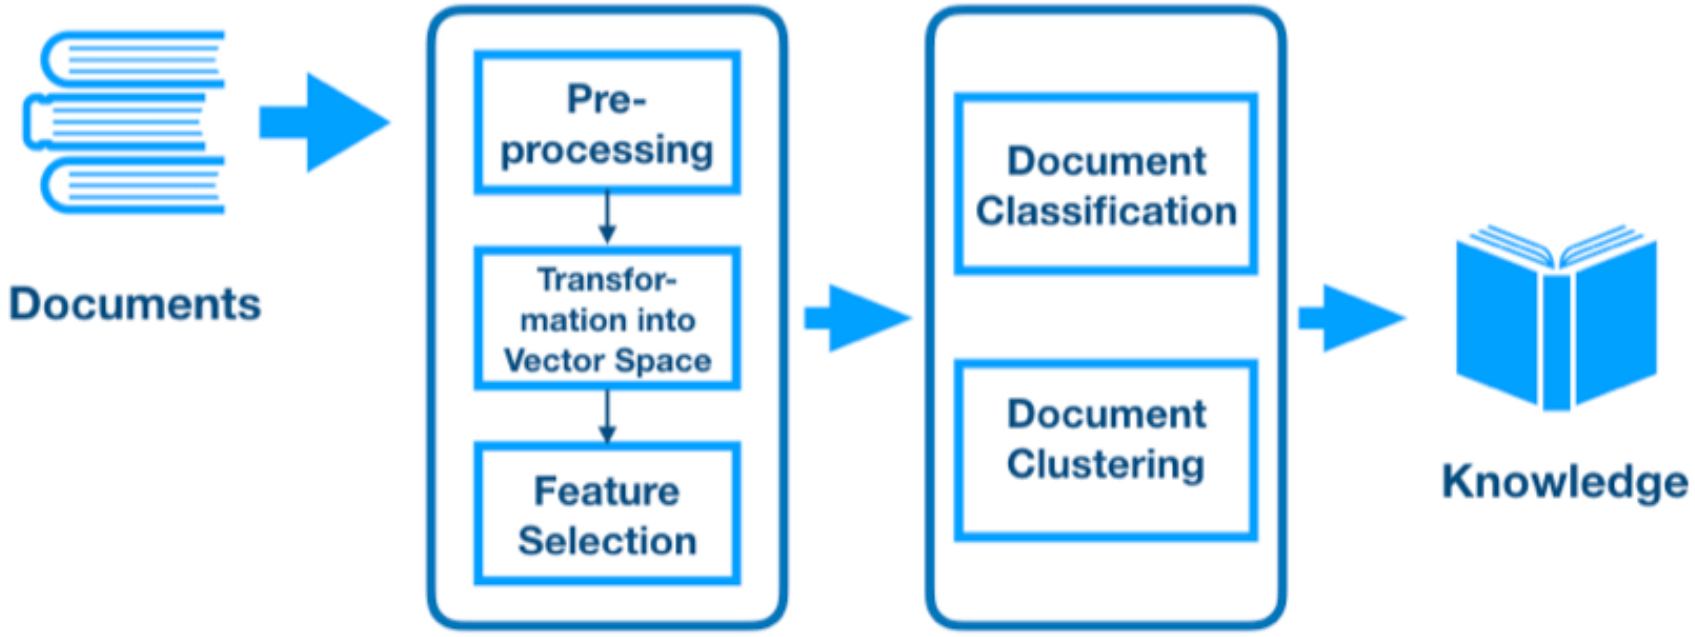

- Text Mining Process Flow

(Source: Miner and Eder et al)

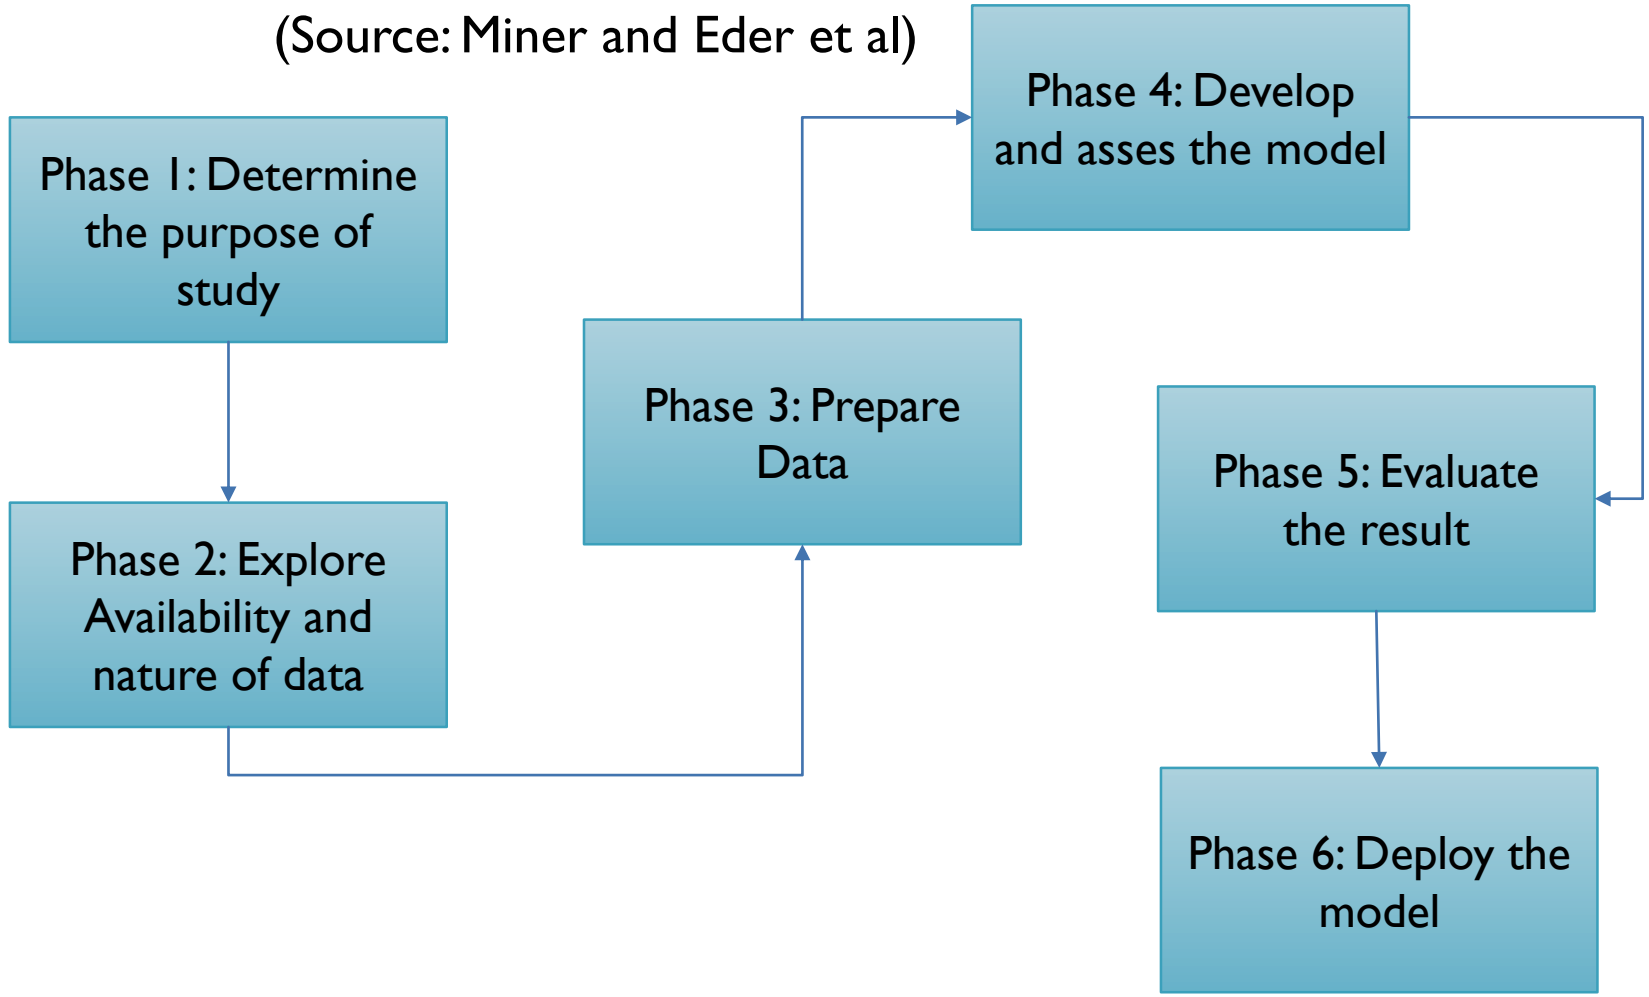

- Text Mining Process Flow

(Source: Data Flair)

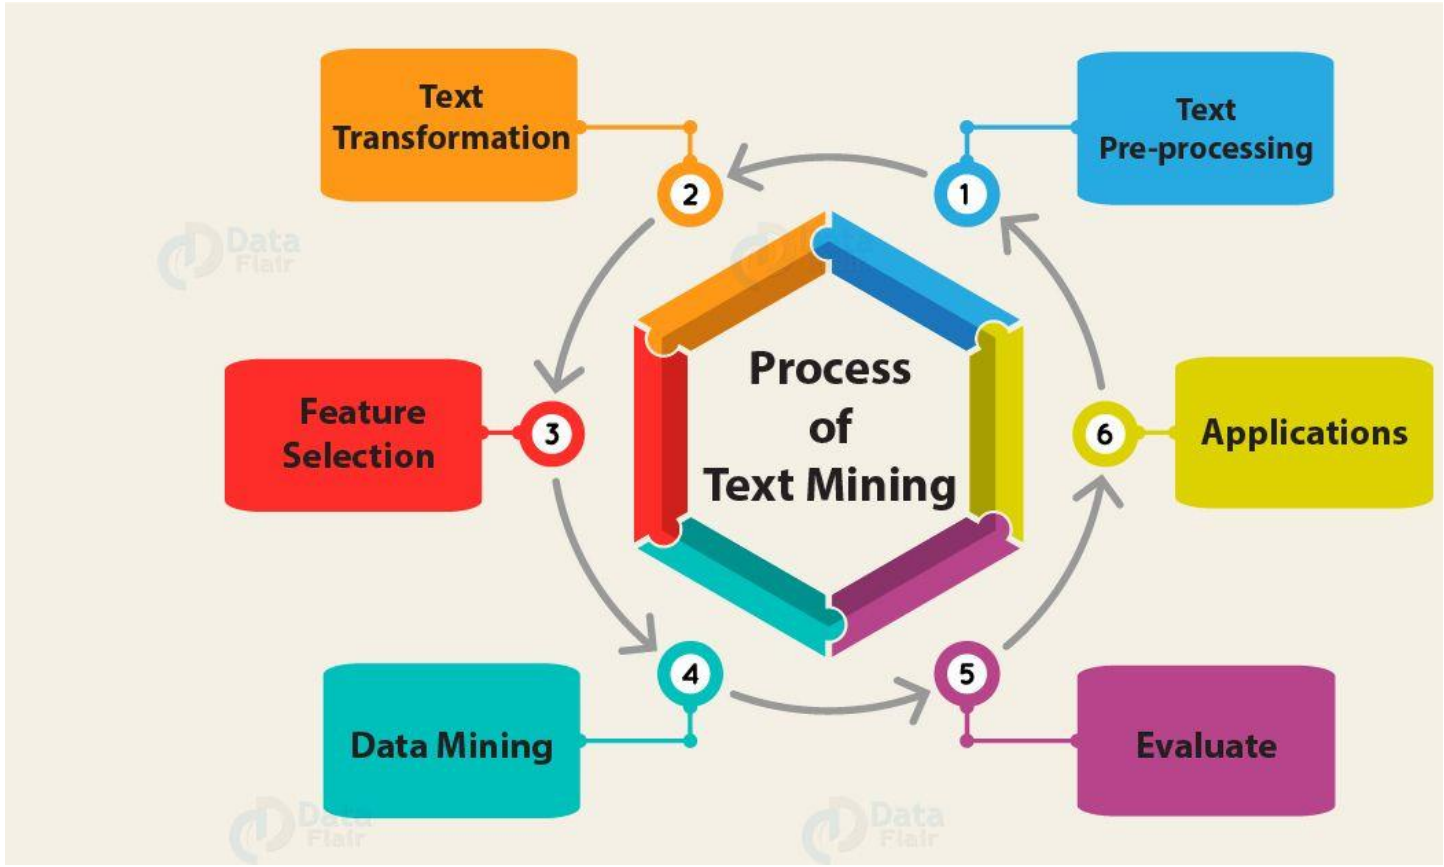

- Text Mining is not easy
  - Language is ambiguous

- **Homonomy**: same word, different meaning by accident of history
  - Bank
    - a. Mary walked along the bank of the river.
    - b. HarborBank is the richest bank in the city.

Synonymy: synonyms, different words, similar or same meaning; can substitute one word for the other without changing the meaning of the sentence substantively.

Synonyms can have differing connotations...

a. Miss Nelson became a kind of big sister to Benjamin.

b. Miss Nelson became a kind of large sister to Benjamin.

Polysemy: same word or form, but different, albeit related meaning

Bank

a. The bank raised its interest rates yesterday.

b. The store is next to the newly constructed bank.

c. The bank appeared first in Italy in the Renaissance.

Hyponymy: concept hierarchy or subclass (subordinates)

Animal (noun)

a. dog

b. cat

Injury

a. Broken leg, contusion...

TEXT MINING

- Concepts and word extraction usually results in a huge dimension
  - Thousands of new fields
  - Each field typically has low information
- Misspelling abbreviations, spelling variants

TEXT MINING

Label	Challenges
Words and morphemes	Word segmentation: dividing text into words. Fairly easy for English and other languages that use whitespace; much harder for languages like Chinese and Japanese. – Assigning part of speech. – Identifying synonyms; synonyms are useful for searching. – Stemming: the process of shortening a word to its base or root form. For example, a simple stemming of *words* is *word*. – Abbreviations, acronyms, and spelling also play important roles in understanding words.

TEXT MINING

Multiword and sentence	Phrase detection: *quick red fox*, *hockey legend Bobby Orr*, and *big brown shoe* are all examples of phrases. – Parsing: breaking sentences down into subject-verb and other relationships often yields useful information about words and their relationships to each other. – Sentence boundary detection is a well-understood problem in English, but is still not perfect. – Coreference resolution: "Jason likes dogs, but he would never buy one." In this example, *he* is a coreference to Jason. The need for coreference resolution can also span sentences.

TEXT MINING

Multiword and sentence	– Words often have multiple meanings; using the context of a sentence or more may help choose the correct word. This process is called *word sense disambiguation* and is difficult to do well. – Combining the definitions of words and their relationships to each other to determine the meaning of a sentence
Multisentence and paragraph	At this level, processing becomes more difficult in an effort to find deeper understanding of an author's intent. Algorithms for summarization often require being able to identify which sentences are more important than others.

OUTLINE

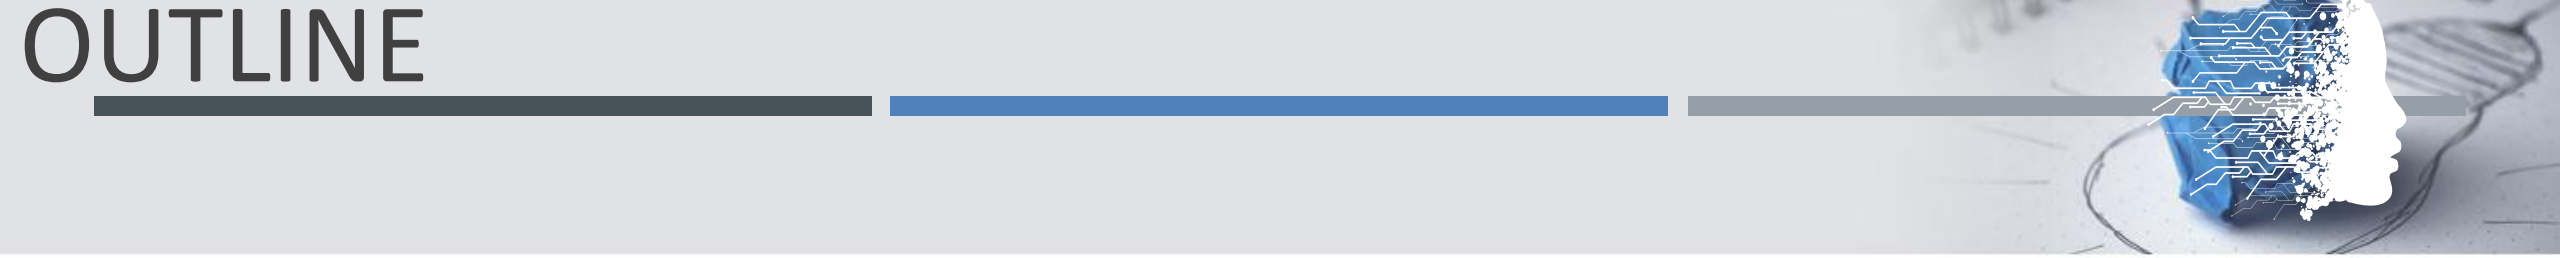

Introduction to Machine Learning

LEARNING – HOW IT IS DONE?

- Mr. Micky found unknown food
- He takes the first bite

Scenario 1: Learning

- Mr. Micky experience poisonous features in food

- Mr. Micky experience no poisonous features in food

- Mr. Micky eats known food without second thought
- He has knowledge (through experience) of the features of the food

Scenario 2: How to

HOW MACHINE LEARNS

- Machine learns through **computer programs** (the language of machine) which learns of from **experience**.

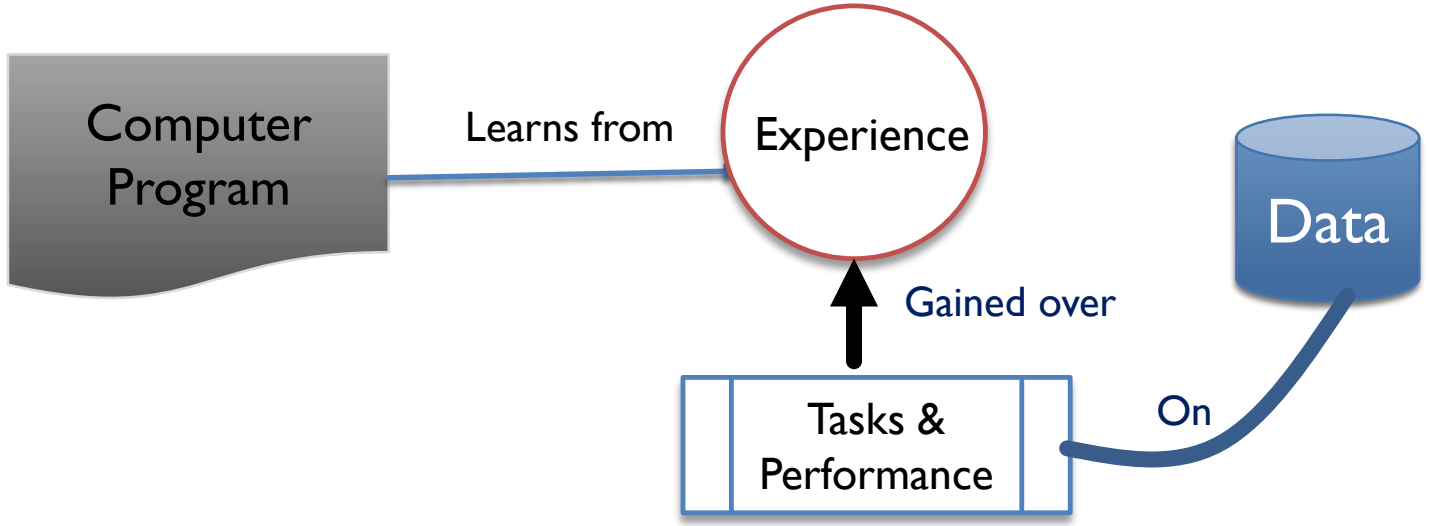

- Machine learns from **changes** in computer program or **data** (input); e.g., samples of speech from different person improve speech recognition.

MACHINE LEARNING APPROACHES

- Definitions of Machine Learning
  - Machine learning is a method that gives computer the ability to learn **without** being explicitly programmed (Arthur, S.)
  - Machine learning is a set of **tools** that, broadly speaking, allow to "**teach**" computers how to perform tasks by providing examples of how they should be done (Mitchel, T.)
  - Machine learning is a method of data analysis that **automates** analytical model building (SAS)

LEARNING METHODS

- Learning methods of machines
  - Inductive learning: It involves the process of learning by example – where a system tries to induce a model (general rule/function) from training data.

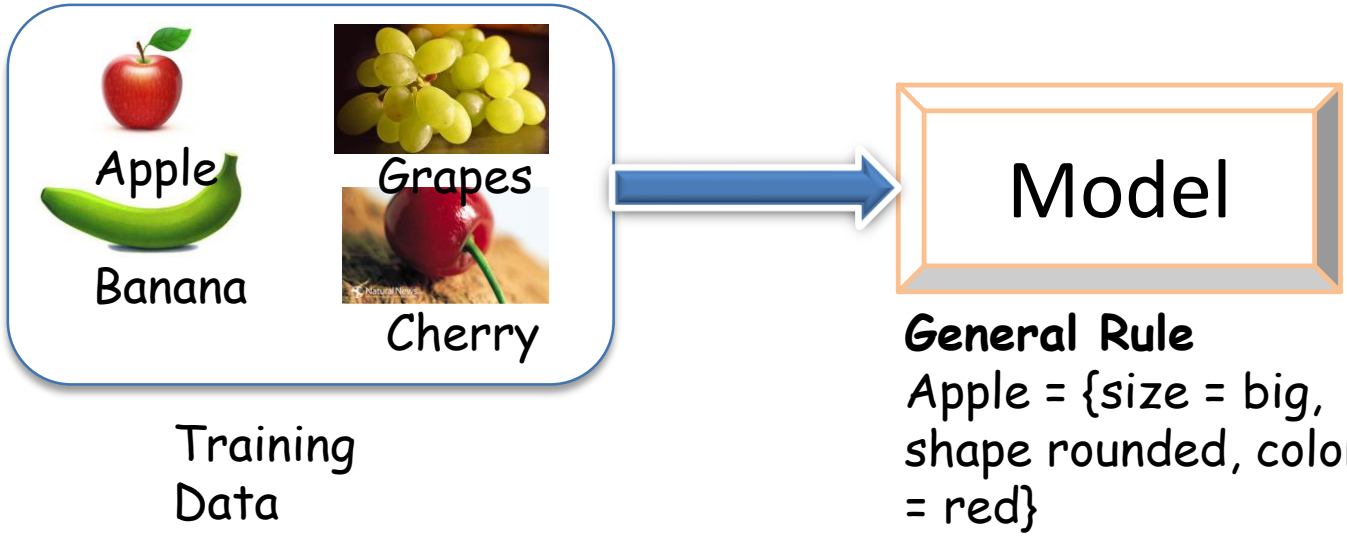

LEARNING METHODS

- Deductive learning: It applies the learning model to predict the outcome.

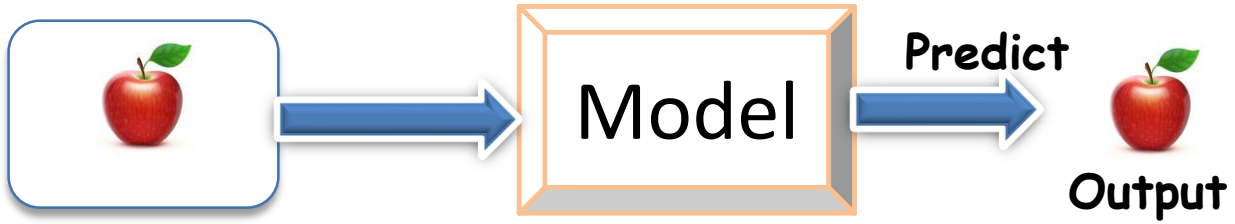

- Instance-based learning
  - The most trivial form of learning
  - The system learns the examples by heart, then generalizes to new cases by using a similarity measure to compare them to the learned examples (or a subset of them)

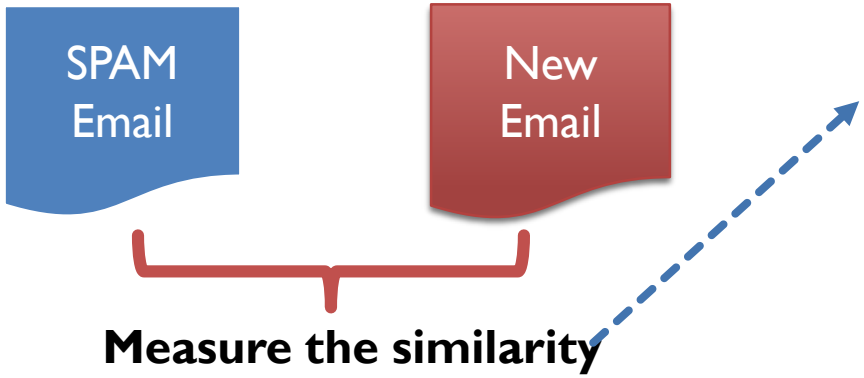

- Find common words between two mails
- If number of common words is significant then flag the new email as Spam

LEARNING METHODS

- Model-based learning
  - Generalize from a set of examples is to build a model of these examples and then use that model to make predictions

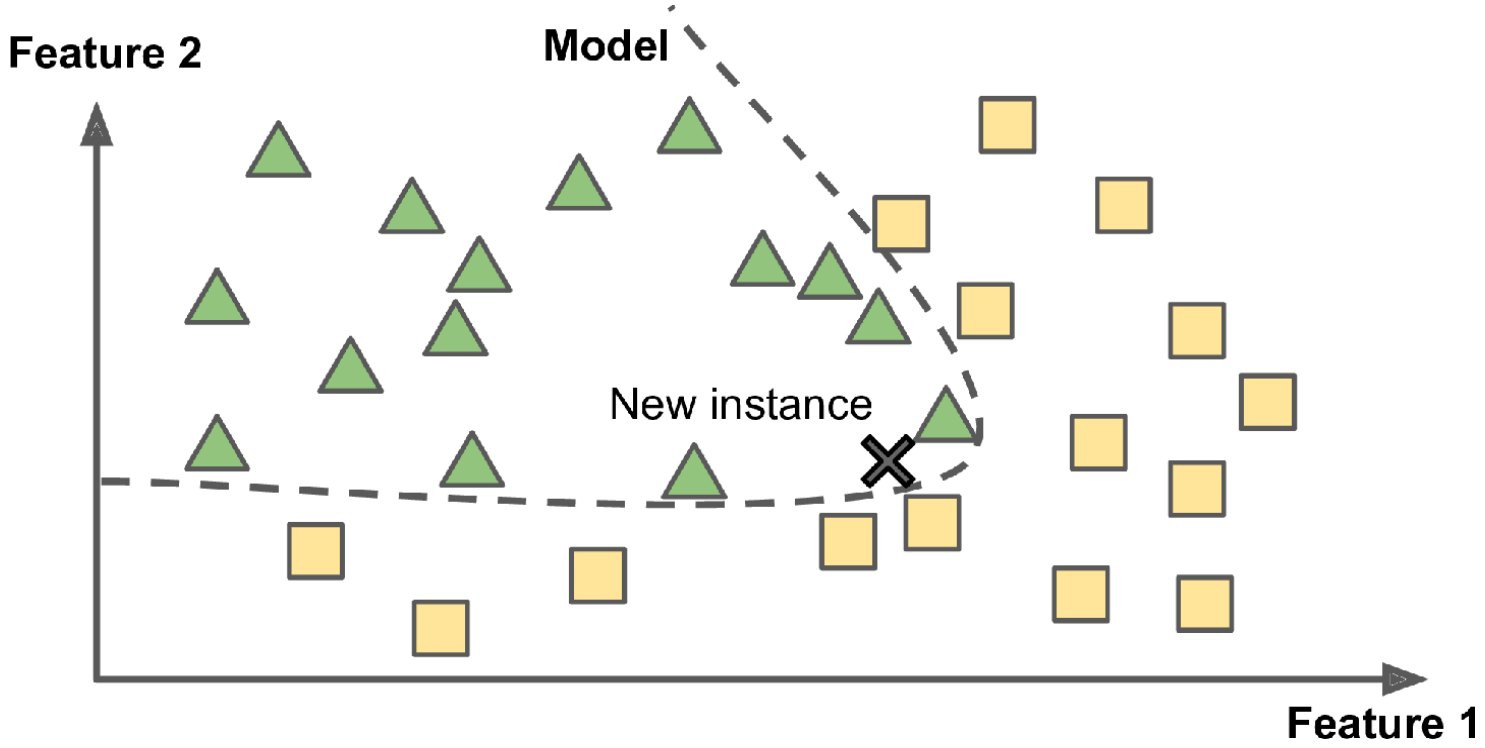

MACHINE LEARNING APPROACHES

- There are different approaches for machine learning systems
  - **Semi Automatic**
  - **Fully Automatic**

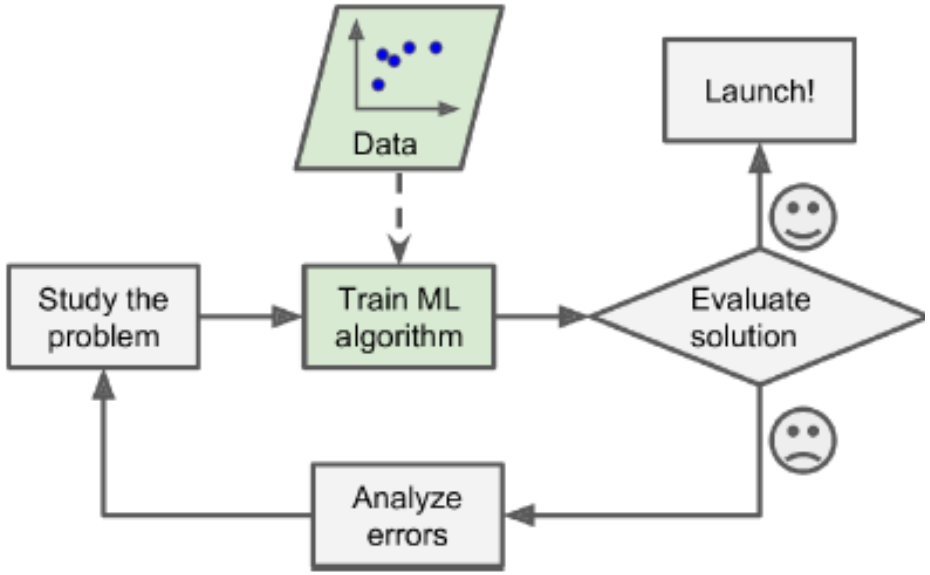

Semi-automatic system approach

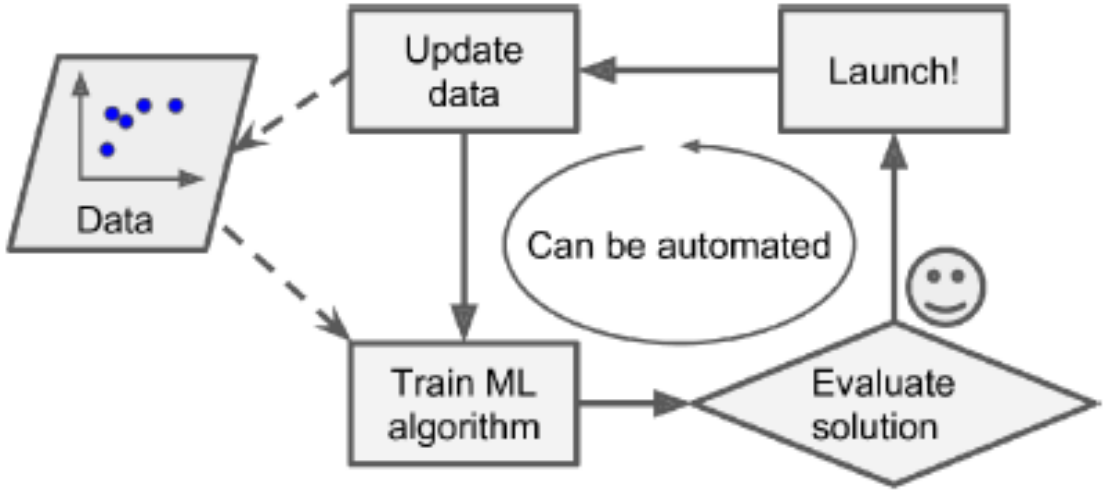

Automatic system approach

MACHINE LEARNING APPROACHES

- Machine learning helps human to learn

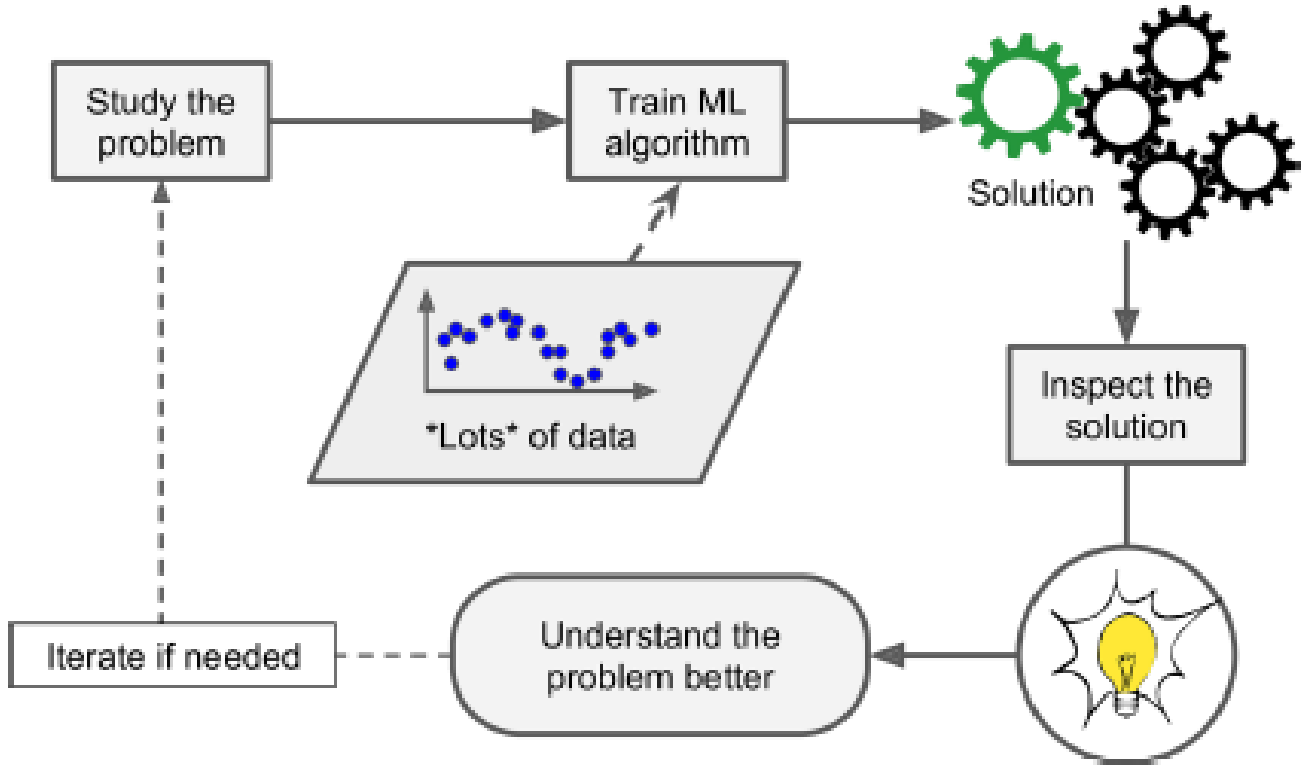

TYPES OF MACHINE LEARNING SYSTEM

- **Supervised Learning**: In supervised learning, the training set you feed to the algorithm includes the desired solutions, called **labels**

Data: $D = \{d_1, d_2, .., d_n\}$ **a set of *n* examples**

$d_i = < \mathbf{x}_i, y_i >$

$\mathbf{x}_i$ is input vector, and y is desired output

Objective: learn the mapping $f : X \to Y$

s.t. $y_i \approx f(x_i)$ for all $i = 1, .., n$

Two types of problems:

- **Regression:** X discrete or continuous → Y is **continuous**
- **Classification:** X discrete or continuous → Y is **discrete**

Ex: Price of a car, given a set of *features*

Ex: Classifying emails along with their class spam or Not Spam

- **Note:** Some regression algorithms can be used for classification as well, and vice versa.
  - For example, Logistic Regression is commonly used for classification, as it can output a value that corresponds to the probability of belonging to a given class

TYPES OF MACHINE LEARNING SYSTEM

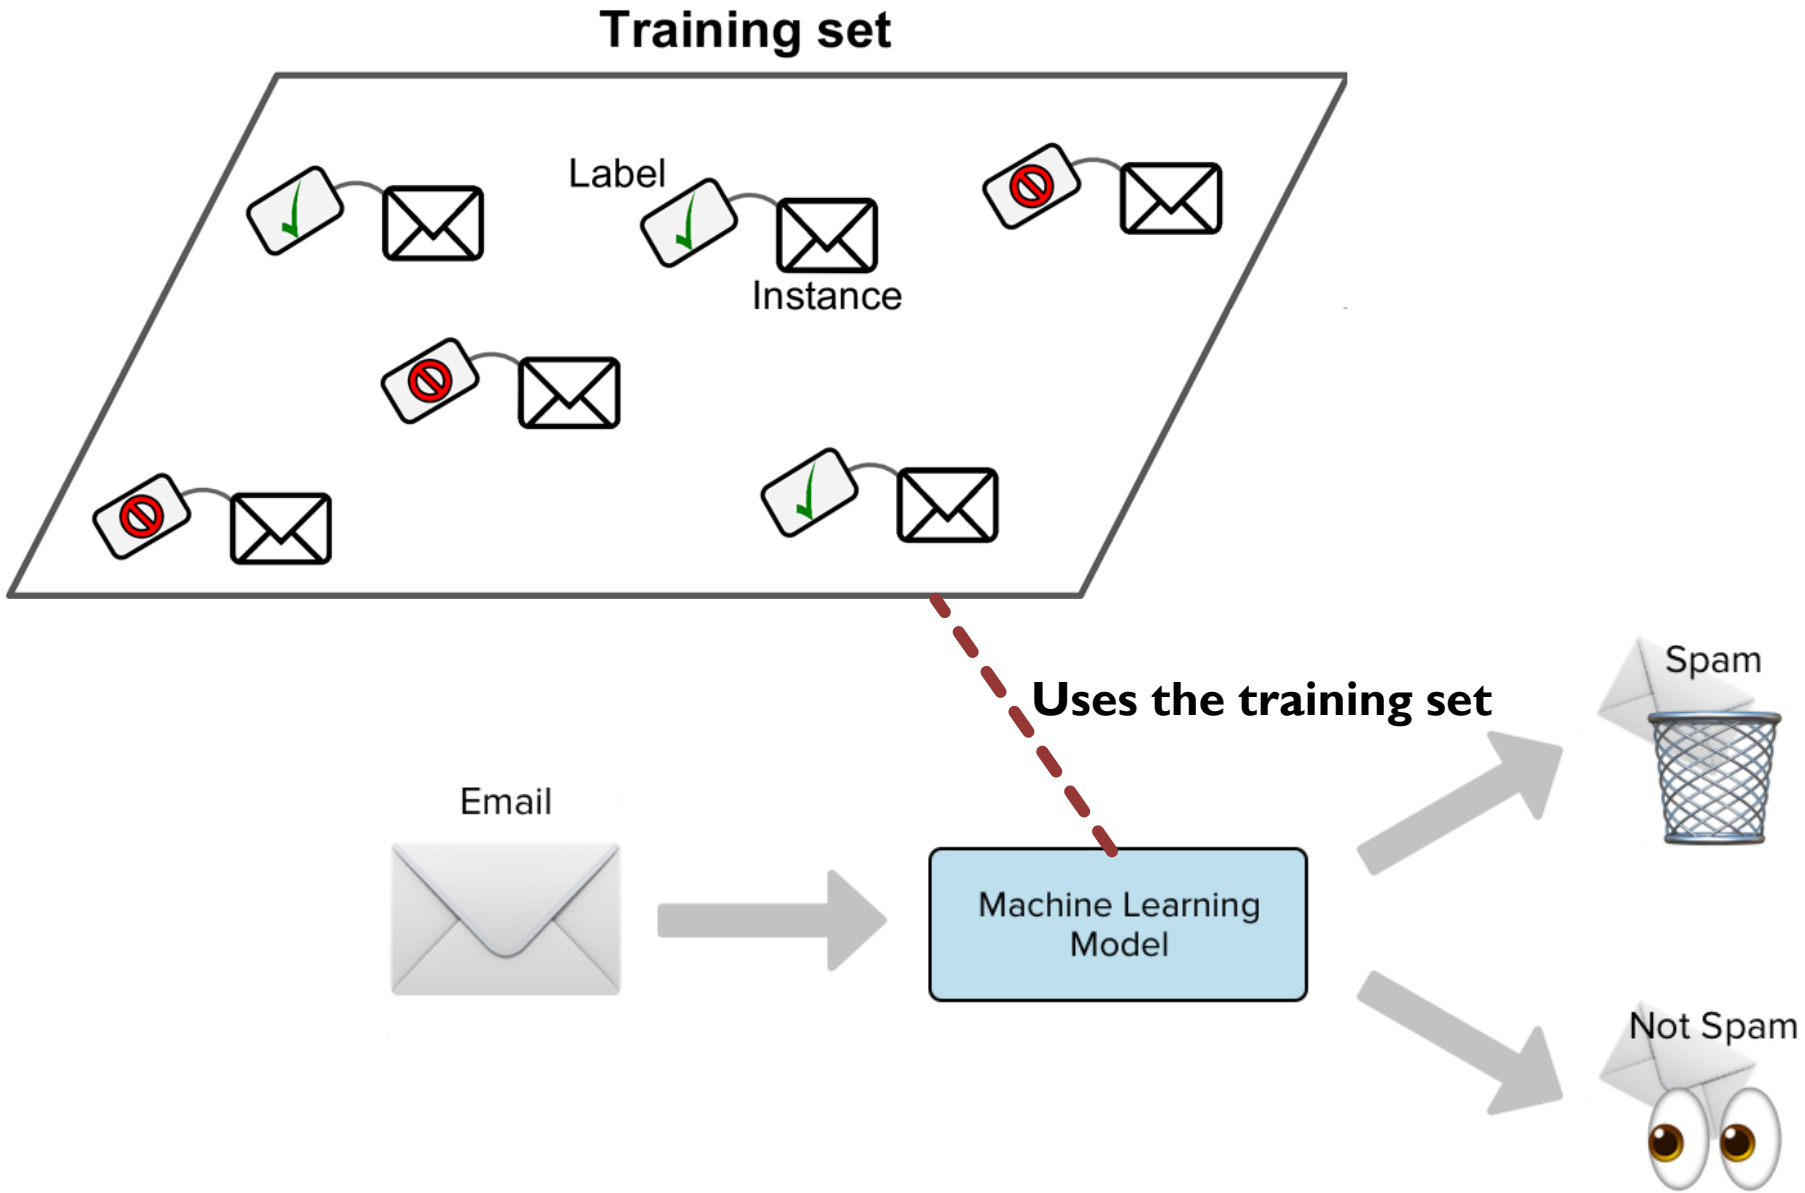

TYPES OF MACHINE LEARNING SYSTEM

- The most widely used Supervised Learning Algorithms
  - k-Nearest Neighbors
  - Linear Regression
  - Logistic Regression
  - Support Vector Machines (SVMs)
  - Decision Trees and Random Forests
  - Neural networks

TYPES OF MACHINE LEARNING SYSTEM

- **Unsupervised Learning**
  - It is a type of learning where algorithm used to draw inferences from datasets consisting of input data (x) without labeled responses.
  - Goal:
    - In some pattern recognition problems, the training data consists of a set of input vectors x without any corresponding target values.
    - The goal in such unsupervised learning problems may be to discover groups of similar examples within the data.

TYPES OF MACHINE LEARNING SYSTEM

- **Unsupervised Learning**
  - Training data provides "example", but we have no specific outcomes.
  - In simple word **there is no label associated with this learning.**
  - In unsupervised learning the machine tries to find interesting patterns in the data.

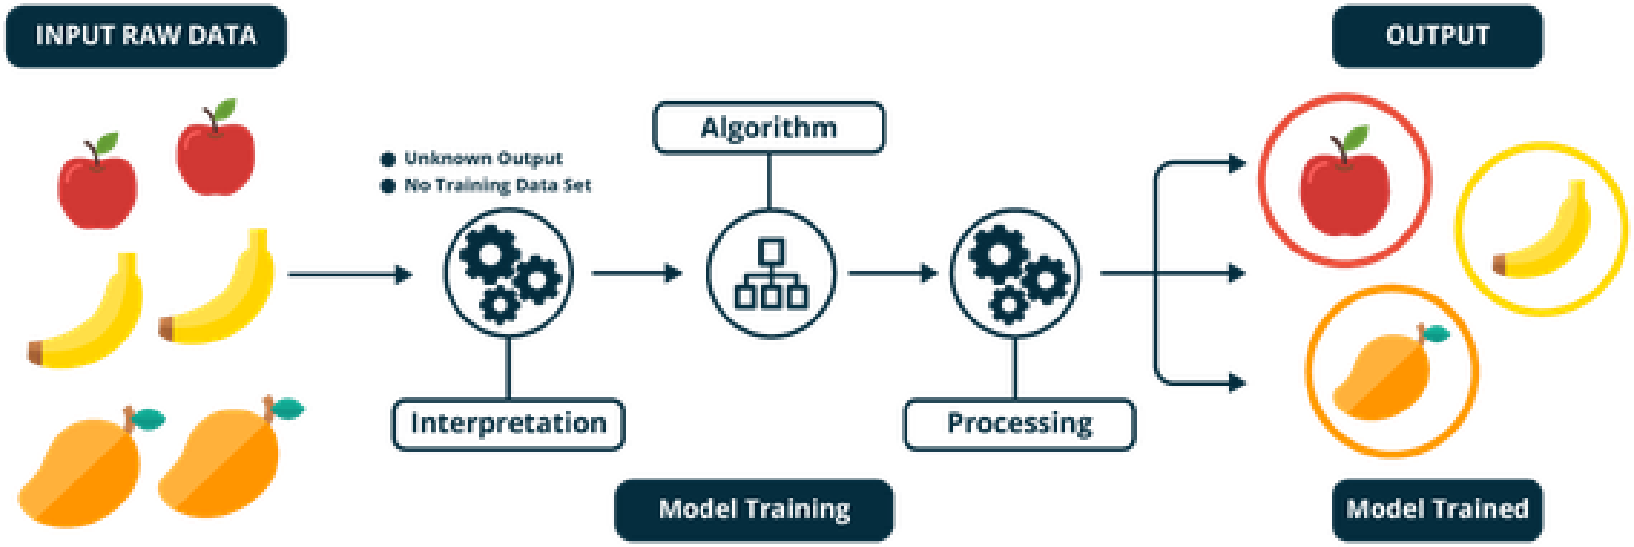

TYPES OF MACHINE LEARNING SYSTEM

- **Unsupervised Algorithms**

Unsupervised Algorithms	
Clustering	—K-Means —DBSCAN —Hierarchical Cluster Analysis (HCA)
Anomaly detection and novelty detection	—One-class SVM —Isolation Forest
Visualization and dimensionality reduction	—Principal Component Analysis (PCA) —Kernel PCA —Locally Linear Embedding (LLE) —t-Distributed Stochastic Neighbor Embedding (t-SNE)
Association rule learning	—Apriori —Eclat

TYPES OF MACHINE LEARNING SYSTEM

- **Semi-supervised Learning**
  - Semi-supervised learning (SSL) is halfway between supervised and unsupervised learning
  - In addition to unlabelled data, the algorithm is provided with some super-vision information – but not necessarily for all examples

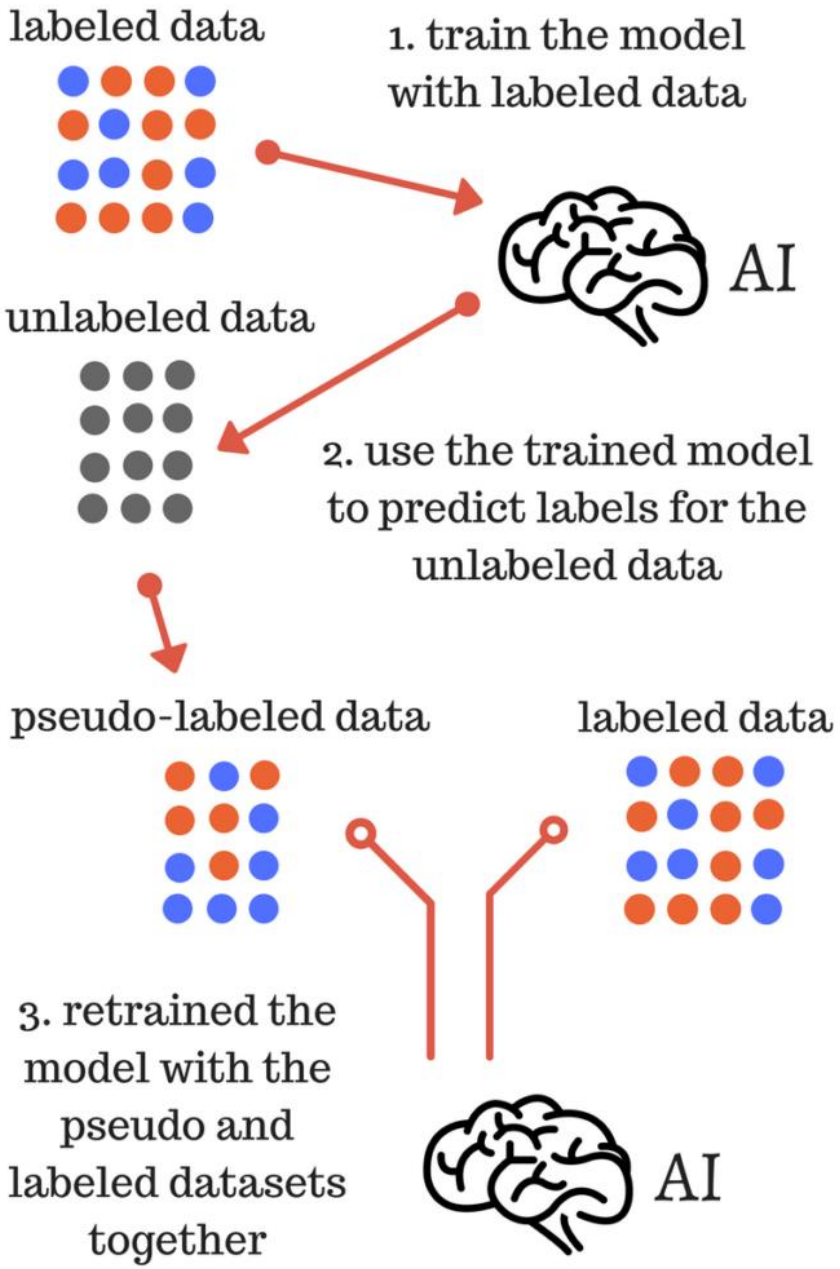

TYPES OF MACHINE LEARNING SYSTEM

- **Reinforcement Learning**
  - The learning system, called an agent in this context, can observe the environment, select and perform actions, and get rewards in return (or penalties in the form of negative rewards.
  - It must then learn by itself what is the best strategy, called a policy, to get the most reward over time
  - A policy defines what action the agent should choose when it is in a given situation.

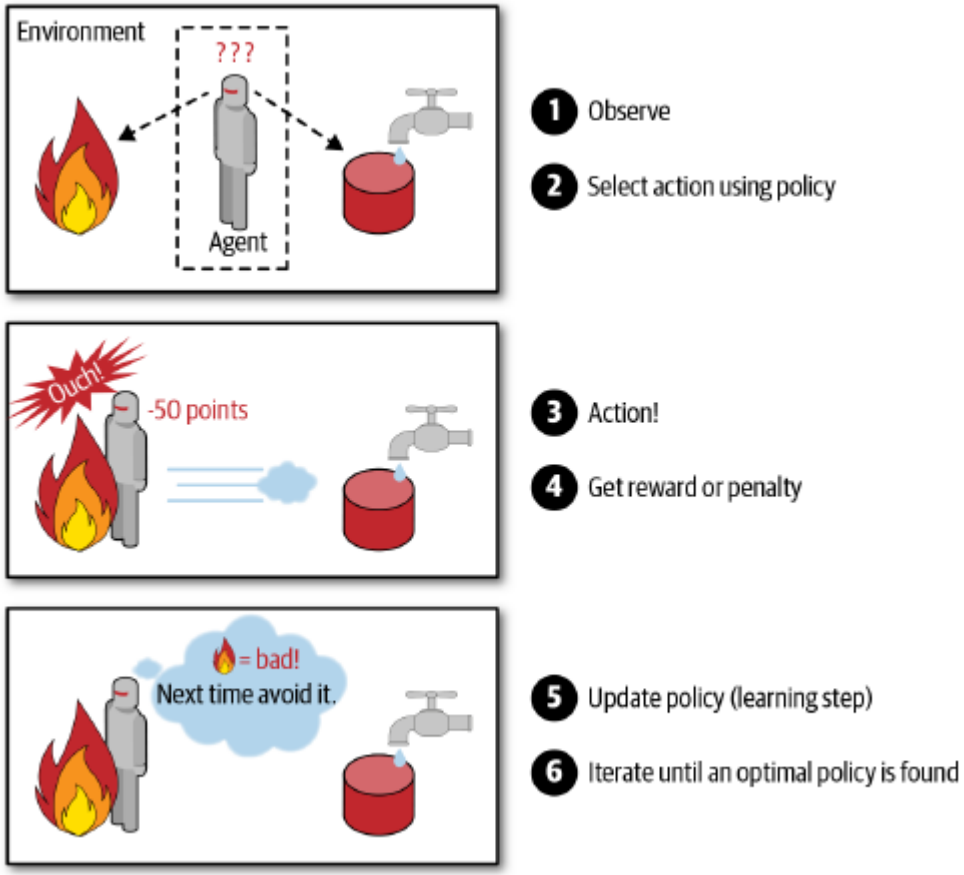

DEPLOYMENT MODELS

- **Batch Learning**
  - The system must be trained using all available data
  - The system is incapable of learning incrementally
  - For new data (such as a new type of spam),
    - Train a new version of the system from scratch on the full dataset (not just the new data, but also the old data)
    - Then stop the old system and replace it with the new one.

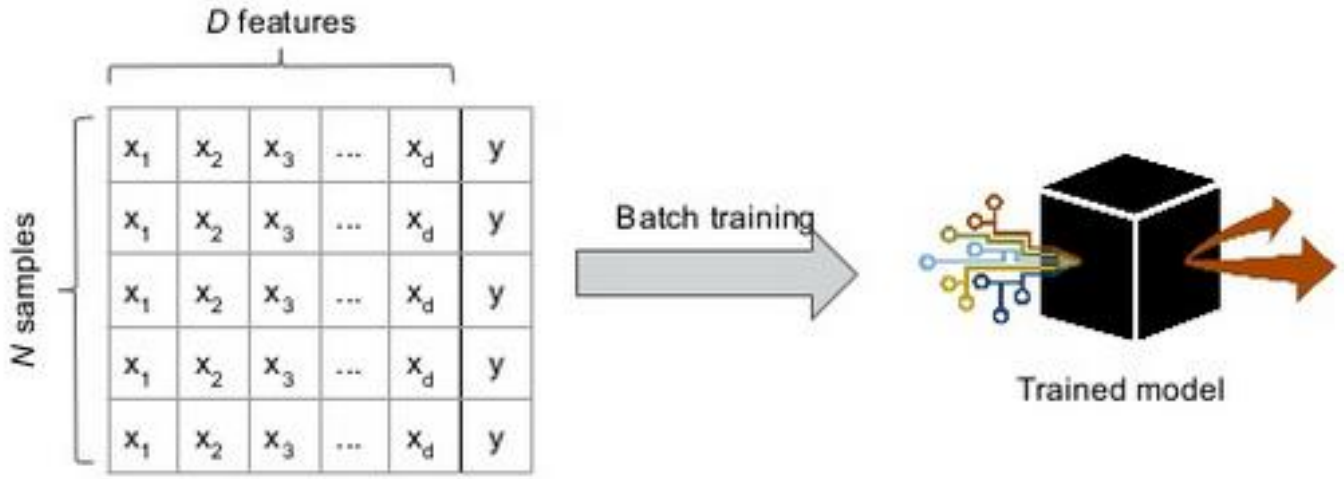

DEPLOYMENT MODELS

- **Online Learning**
  - Train the system incrementally by feeding it data instances sequentially, either individually or in small groups called mini-batches.
  - Each learning step is fast and cheap, so the system can learn about new data on the fly, as it arrives

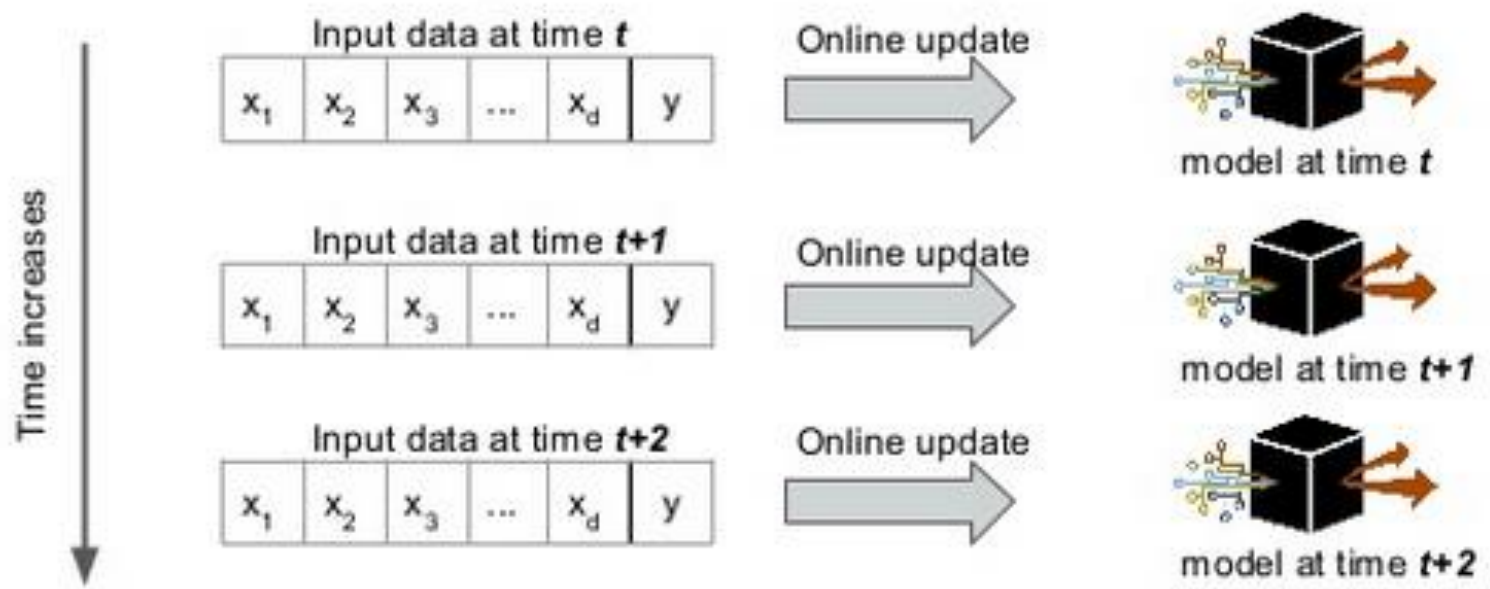

CHALLENGES OF MACHINE LEARNING

- **Insufficient Quantity of Data**
  - For very simple problems you typically need thousands of examples
  - For complex problems such as image or speech recognition you may need millions of examples

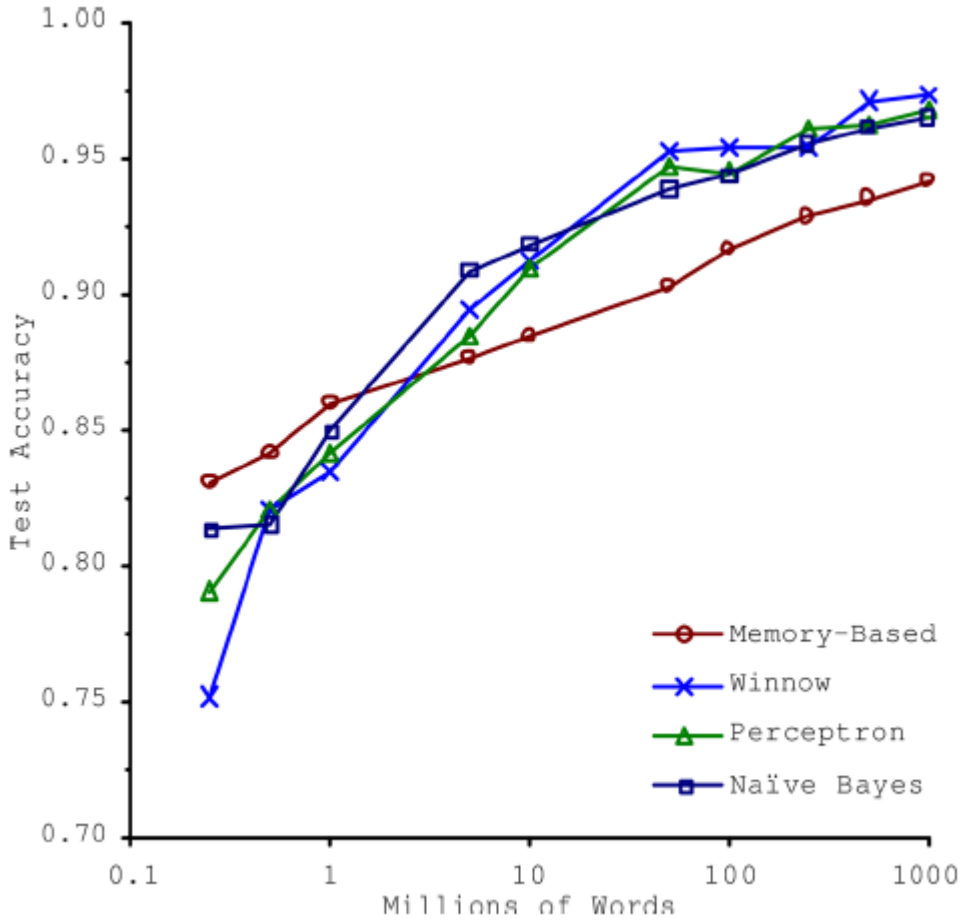

CHALLENGES OF MACHINE LEARNING

- **Non-Representative Training Data**
  - It is crucial that your training data be representative of the new cases you want to generalize to

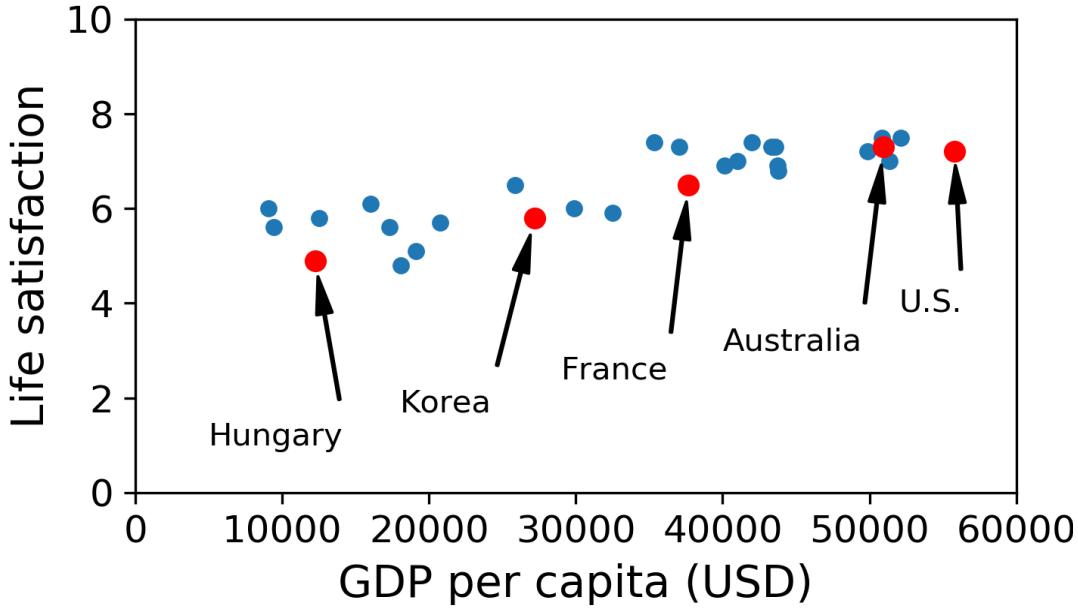

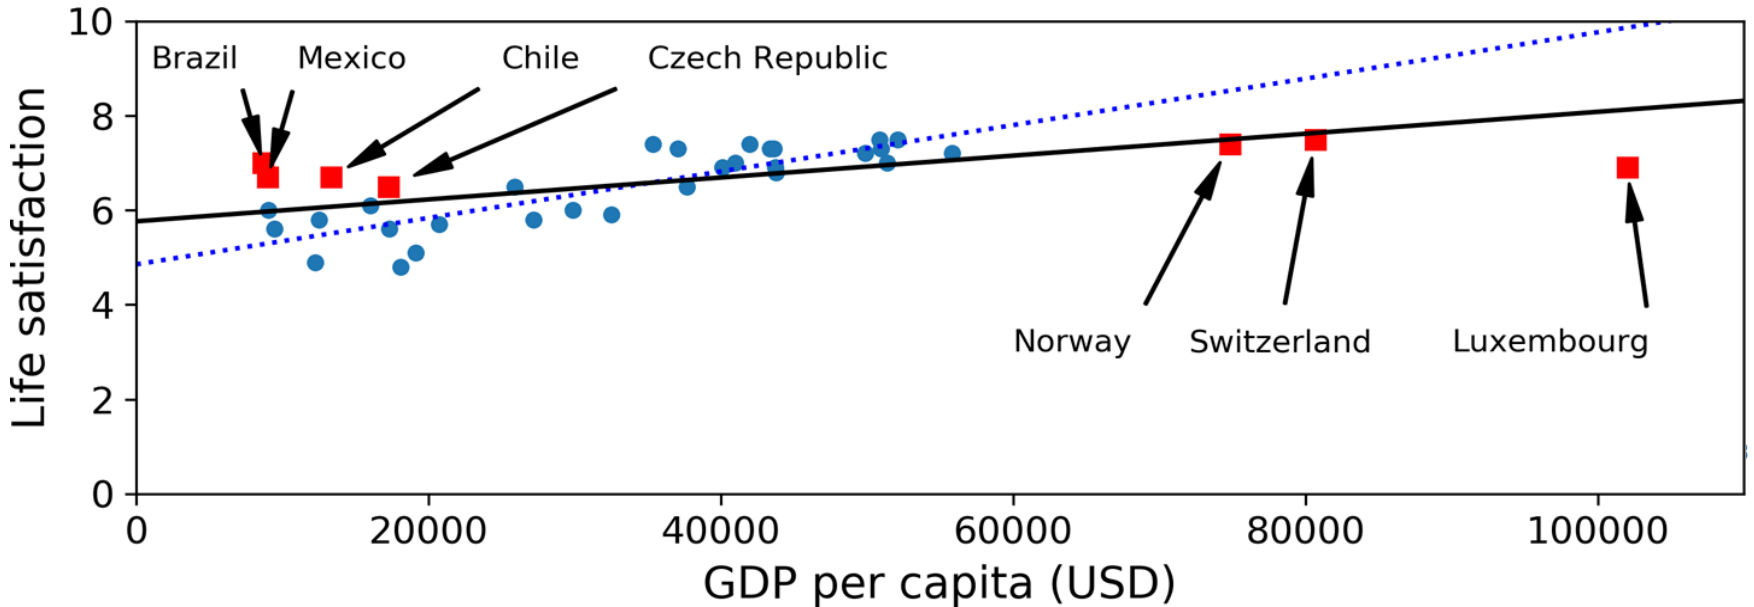

CHALLENGES OF MACHINE LEARNING

- **Poor Quality Data**
  - If training data is full of errors, outliers, and noise it will make it harder for the system to detect the underlying patterns, so your system is less likely to perform well
- **Overfitting training data**
  - The model performs well on the training data, but it does not generalize well.
- **Underfitting**
  - It occurs when your model is too simple to learn the underlying structure of the data

CHALLENGES OF MACHINE LEARNING

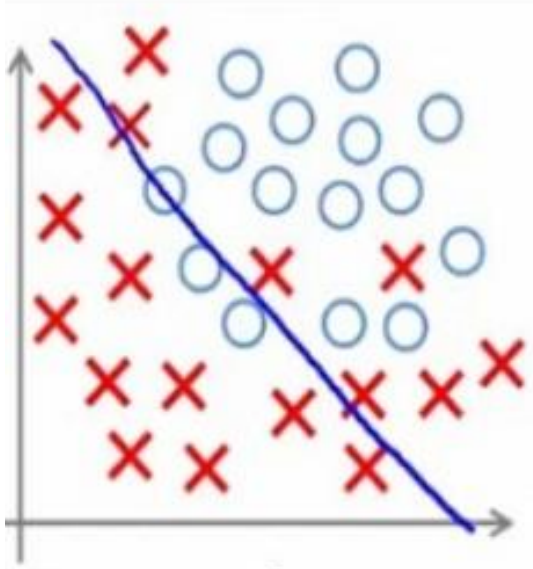

Under-fitting

(too simple to explain the variance)

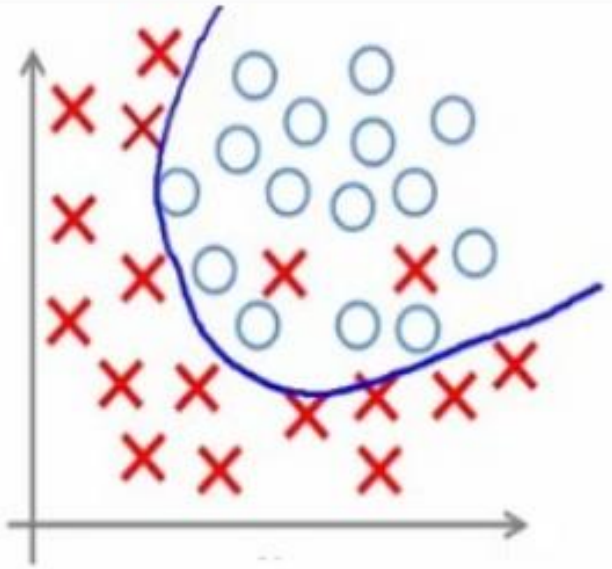

Appropriate-fitting

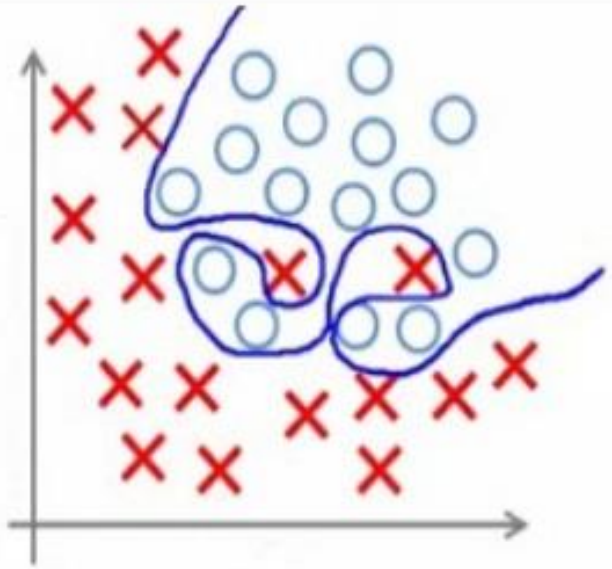

Over-fitting

(forcefitting -- too good to be true)

WHAT PROBLEM WE WILL SOLVE

- **Data Pre-processing**

  Feature extraction

  - Tokenization
  - Removing unnecessary punctuation, tags
  - Removing stop words
  - Stemming
  - Lemmatization/POS tagging

  Feature selection

  - Term frequency
  - Chi-square
  - Expected cross entropy
  - Odds Ratio

WHAT PROBLEM WE WILL SOLVE

- Algorithms
  - Support Vector Machines
  - Naive Bayes
  - Maximum Entropy
  - K-nearest neighbour
  - Decision Tree
  - TF-IDF
  - Conditional Random Field (CRF)
  - Latent Dirichlet Allocation (LDA)
  - Artificial Neural Networks
  - LSTM
  - Bi-directional LSTM
  - Recurrent neural networks
  - Models of word embedding:
    - Word2Vec
    - Glove

WHAT PROBLEM WE WILL SOLVE

- Common NLP techniques
  - Named Entity Recognition
  - Sentiment Analysis
  - Language detection
  - Topic Modeling
  - Text Summarization
  - Machine Translation